

Towards Automatic Lock Removal for Scalable Synchronization

Maya Arbel^{*†}

Guy Golan Gueta^{*}

Eshcar Hillel^{*}

Idit Keidar^{*†}

Yahoo Labs, Haifa, Israel

The Technion, Haifa, Israel

ABSTRACT

We present a *code transformation* for concurrent data structures, which increases their scalability without sacrificing correctness. Our transformation takes lock-based code, and replaces some of the locking steps therein with optimistic synchronization, in order to reduce contention. The main idea is to have each operation perform an optimistic traversal of the data structure as long as no shared memory locations are updated, and then proceed with pessimistic code. The transformed code inherits essential properties of the original one, including linearizability, serializability, and deadlock freedom.

Our work complements existing pessimistic transformations that make sequential code thread-safe by adding locks. In essence, we provide a way to optimize such transformations by reducing synchronization bottlenecks (for example, locking the root of a tree). The resulting code scales well and significantly outperforms pessimistic approaches. We further compare our synthesized code to state-of-the-art data structures implemented by experts. We find that its performance is comparable to that achieved by the custom-tailored implementations. Our work thus shows the promise that automated approaches bear for overcoming the difficulty involved in manually hand-crafting concurrent data structures.

1. INTRODUCTION

1.1 Generic Lock Removal

The steady increase in the number of cores in today’s computers is driving software developers to allow more and more parallelism. An important focal point for such efforts is scaling the concurrency of shared data structures, which are often a principal friction point among threads. Many recent works have been dedicated to developing scalable concurrent data structures (e.g., [8, 18, 40, 12, 15, 9, 5, 20, 11, 29, 23, 38]), some of which are widely used in real-world systems [44].

Each of these projects generally focuses on a single data structure (for example, a binary search tree [11] or a queue [38]) and manually optimizes its implementation. These data structures are developed by concurrency experts, typically PhDs or PhD candidates. Proving the correctness of such custom-tailored data structures is painstaking; for example, the proofs of [9, 20] are 31 and 20 pages long respectively. The rationale behind dedicating so much effort to one data structure is that it is generic and can be used by many applications. Nevertheless, systems often use data struc-

tures in unique ways that necessitate changing or extending their code (e.g., [2, 3, 43, 46]), which limits the usability of custom-tailored implementations. Hence, the return-on-investment for such endeavors may be suboptimal. Here, we propose an approach to facilitate this labor-intensive process, making scalable synchronization more readily available.

Specifically, we present in Section 2 an algorithm for a source-to-source code transformation that takes a lock-based concurrent data structure implementation as its input and generates more scalable code for the same data structure via judicious use of optimism. Our approach combines optimism and pessimism in a new, practical, way. In striking the balance between the two, we exploit the common access pattern in data structure operations, (for example, tree insertion or removal), which typically begin by traversing the data structure (to the insertion or removal point), and then perform (mostly) local updates at that location. Our transformation replaces locking steps in the initial read-only traversal of each operation with optimistic synchronization, whereas the update phase employs the original lock-based synchronization. Our work may thus be seen as a form of software lock elision for read-only prefixes of operations (transactions).

Combining optimism and pessimism allows us to achieve “the best of both worlds” – while the optimistic traversal increases concurrency and eliminates bottlenecks, the use of pessimistic updates saves the overhead associated with speculative or deferred shared memory updates, (as occurs in software transactional memory [27]). The partially-optimistic execution is compatible with the original code, which permits us to re-execute operations pessimistically when too many conflicts occur, avoiding livelocks.

Moreover, our transformation refrains from introducing a shared global clock (as used in some transactional memory systems [45]) or other sources of contention. Thus, if the original code is *disjoint access parallel* [32], i.e., threads that access disjoint (abstract) data objects do not contend on (low level) shared memory locations, then this property holds also for the transformed code.

We show in Section 3 that our transformation preserves the external behavior of the original lock-based code; formal proofs are deferred to Appendix A. In other words, if the original code is correct (in the sense of serializability, linearizability, and deadlock-freedom), so is the transformed version.

1.2 Towards Fully Automatic Parallelization

One important use case for our transformation is to apply it in conjunction with automatic lock-based parallelization mechanisms [24, 36]. The latter instrument sequen-

tial code and add fine-grained lock and unlock instructions that ensure its safety in concurrent executions. Our evaluation shows that, by themselves, solutions of this sort may scale poorly. This is due to synchronization bottlenecks, e.g., the root of a tree, which is locked by all operations. By subsequently applying our transformation, one can optimize the lock-based code they produce, yielding an end-to-end approach to scalable parallelization of sequential code. Furthermore, our transformation, as well as these locking-based parallelization mechanisms, appear to be amenable to compile-time implementation, and can thus potentially lead to the development of automatic tools for efficient parallelization of sequential code, which is beyond the scope of the current paper.

1.3 Evaluation

In Section 4 we evaluate our transformation by generating three data structures— an unbalanced search tree, a treap (randomized balanced search tree), and a skip list that supports range queries. We synthesize the first two from sequential implementations using the algorithm of [24] (*domination locking*), followed by our transformation. For the skip list, we manually add fine-grained locks (in a straightforward manner), and then apply our transformation. All examples are implemented in Java. We evaluate the scalability of the resulting code in a range of workload scenarios on a 32-core machine. In all cases, the lock-based implementations do not scale – their throughput remains flat as the number of running threads increases. In contrast, the code generated by our transformation is scalable, and its throughput continues to grow with the number of threads.

We compare our synthesized code to state-of-the-art data structures that were hand-crafted by experts in the field [18, 11, 1, 23]. Our results show that the search tree and treap implementations we have generated perform comparably to custom-tailored solutions.

We further consider a data structure that supports range queries, which are required by many applications (e.g., [2, 22]). To this end we implement a skip list – a data structure that naturally supports range queries. While range queries implemented using the iterators available in the Java concurrency library’s skip list [1] perform somewhat better than ones in our synthesized code, it is important to note that these iterators are *not* linearizable (atomic), and only support so-called weak consistency, whereas range queries in our implementation are linearizable. A notable previous linearizable implementation is due to Bronson et al. [11], and it is significantly out-performed by our synthesized code. This illustrates the benefit of the broad applicability of our generic approach compared to specific custom-tailored implementations.

Note that the alternative generic approach we compare with is domination locking [24]. Software transactional memory is another general-purpose approach for parallelization, yet due to the significant overhead associated with this approach [13, 19], it is often complemented by custom-tailored data structure implementations whose operations can be called from within transactions [28, 34, 10]. Further discussion of related work appears in Section 5.

To conclude, this paper demonstrates that generic synchronization, based on a careful combination of optimistic and pessimistic concurrency control, is a promising approach for bringing legacy code to emerging computer architectures.

While this paper illustrates the method for tree and skip list data structures, we believe that the general direction may be more broadly applicable, and maybe used with a variety of locking schemes, such as two phase locking. Section 6 concludes the paper and touches on some directions for future work.

2. TRANSFORMATION

We present an algorithm for a source-to-source transformation, whose goal is to optimize the code of a given data structure implemented using lock-based concurrency control. In Section 2.1, we detail our assumptions about the given code and the locks it uses.

The transformation produces a combination of pessimistic and optimistic synchronization by replacing part of the locking (in the given code) with optimistic concurrency control. Section 2.2 overviews our general approach to combining optimism and pessimism.

The synthesized code consists of three phases - an optimistic read-only phase, a validation phase, and a pessimistic update phase. Section 2.3 details how the code of each of the three phases is produced.

2.1 Lock-based Data Structures

A *data structure* defines a set of *operations* that may be invoked by clients of the data structure, potentially concurrently.

Operations have parameters and local variables, which are private to the invocation, and are thus thread-local. The operations interact via *shared memory variables*, which are also called *shared objects*. Each shared object supports atomic *read* (load) and *write* (store) instructions. More formal definitions of the model appear in Appendix A.

In addition, each shared object is associated with a unique lock, and supports atomic *lock*, *tryLock*, *unlock* and *isLockedByAnotherThread* instructions. Locks are exclusive (i.e., a lock can be held by at most one thread at a time). The execution of a thread trying to acquire a lock (by a *lock* instruction) which is held by another thread is blocked until a time when the lock is available (i.e., is not held by any thread). The other instructions never block. The *tryLock* instruction returns *false* if the lock is currently held by another thread, otherwise it acquires the lock and returns *true*. The *isLockedByAnotherThread* instruction returns *true*, if and only if, the lock is currently held by another thread.

We assume that in the given code every (read or write) access by an operation to a shared object is performed when the executing thread holds the lock on that object. We also assume that the locking in the given code only uses the *lock* and *unlock* instructions.

2.2 Combining Optimism and Pessimism

Generally speaking, optimistic concurrency control is a form of lock-free synchronization, which accesses shared variables without locks in the hope that they will not be modified by others before the end of the operation (or more generally, the transaction). To verify the latter, optimistic concurrency control relies on *validation*, which is typically implemented using version numbers. If validation fails, the operation restarts. Optimistic execution of update operations requires either performing roll-back (reverting variables to their old values) upon validation failure, or deferring writes to commit time; both approaches induce significant

overhead [13]. We therefore refrain from speculative shared memory updates.

The main idea behind our approach is judicious use of optimistic synchronization for reading shared variables without locks, but only as long as the operation does not update shared state. Once an operation writes to shared memory, we revert to pessimistic (lock-based) synchronization. In other words, we rely on validation at the end of the read-only prefix of an operation in order to render redundant locks that would have been acquired and freed before the first update. This scheme is particularly suitable for data structures, since the common behavior of their operations is to first traverse the data structure, and then perform modifications.

Conceptually, our approach divides an operation into three phases: an optimistic *read-only phase*, a pessimistic *update phase* and a *validation phase* that conjoins them. The read-only phase traverses the data structure without taking any locks, while maintaining in thread-local variables sufficient information to later ensure the correctness of the traversal. The read phase is *invisible* to other threads, as it updates no shared variables. The update phase uses the original pessimistic (lock-based) synchronization. The validation phase bridges between the optimistic and pessimistic ones. It first locks the objects for which a lock would have been held at this point by the original locking code, and then validates the correctness of the read-only phase. This allows the update phase to run as if an execution of the original pessimistic synchronization took place. If the validation fails, the operation restarts. In order to avoid livelock, we set a threshold on the number of restarts. If the threshold is exceeded, the code falls back on pessimistic execution. We show below that it is safe to do so, since our semi-optimistic code is compatible with the fully pessimistic one.

Phase Transition.

In many cases, the transition from the read-only phase to the update phase in the original code occurs at a statically-defined code location. For example, many data structure operations begin with a read-only traversal to locate the key of interest, and when it is found, proceed to execute code that modifies the data structure. This is the case in all the examples we consider in Section 4 below.

More generally, it is possible to switch from the optimistic read-only phase (via the validation phase) to pessimistic execution at any point before the first update. Moreover, the phase transition point can be determined dynamically at run time. To this end, the execution mode is tracked with a flag **opt**, initialized to true, indicating the optimistic phase. Every shared memory update operation is instrumented with code that checks **opt**, and if it is true, executes the validation phase followed by setting **opt** to false and continuing the execution from the same location. Every shared memory read operation is instrumented with a check of **opt** to determine whether it is to be executed optimistically or pessimistically.

2.3 Transforming the Code Phases

We now describe how we synthesize the code for each of the phases. The regular three-phase flow is described in Section 2.3.1, and exceptions are described in Section 2.3.2.

2.3.1 Normal Flow

Original Instruction	Transformed Code
<code>x.lock()</code>	<code>if !track(x) then goto S</code>
<code>x.unlock()</code>	<code>lockedSet.remove(x)</code>

Table 1: Transformation for read-only phase: each locking instruction (left side) is replaced with the corresponding code on the right; *S* denotes the beginning of the operation.

We illustrate the transformation for a simple code snippet that adds a new element as the third node in a linked list. The original and transformed code are provided in Figure 1. The latter uses the tracking and validation functions in Figures 2 and 3, resp. For clarity of exposition, we present a statically instrumented version, without tracking the phases using **opt**.

We use *version numbers* to validate the correctness of the optimistic execution of the read-only phase. Our transformation instruments each object *o* with an additional field *version*. If *o* is not locked, then this field represents the current version number of *o*. Note that each object has its own version — i.e., version numbers of different objects are independent of each other.

Read-only Phase.

In this phase, we replace all the lock and unlock instructions with local tracking. In this phase the executing thread is invisible to other thread, i.e., avoids any contention on shared memory both in terms of writing and in terms of locking. During this phase, our synchronization maintains two thread-local multi-sets: *lockedSet* and *readSet*. The *lockedSet* tracks the objects that were supposed to be locked by the original synchronization. The *readSet* tracks versions of all objects read by the operation, in order to allow us to later validate that the operation has observed a consistent view of shared memory.

At the beginning of the read-only phase, we insert code that initializes *lockedSet* and *readSet* to be empty (see line 2 of Figure 1b). We replace every lock and unlock instruction with the corresponding code in Table 1. A lock instruction on object *o* is replaced with code that tracks the object and its version in *lockedSet* and *readSet* (see Figure 2). An unlock instruction on object *o* is replaced with code that removes *o* from *lockedSet*. An example for a transformed code is shown in lines 2-12 of Figure 1b.

In Figure 2 (lines 5-8), we use an eager validation scheme¹: If the object already exists in *readSet*, we check that its current version is equal to the version in *readSet*; and if the versions are different the operation restarts from the beginning (line 5). Similarly, it is also checked to be unlocked; if it is locked, the operation restarts from the beginning (line 7).

[Idit: Need to add discussion of fences and fences in the code.]

Validation Phase.

The code of the validation phase is invoked between the read-only phase and the update phase. This code is shown in lines 13-19 of Figure 1b. It locks the objects that are left in *lockedSet* and validates the objects in *readSet*. To avoid deadlocks, the locks are acquired using a tryLock in-

¹This validation scheme may be omitted, since it is not required for correctness.

<pre> 1: FUNCTION addThird(List list, Node new) ----- > read-only phase 2: 3: list.lock() 4: Node prev = list.head 5: prev.lock() 6: list.unlock() 7: Node succ = prev.next 8: succ.lock() 9: prev.unlock() 10: prev = succ 11: succ = succ.next 12: succ.lock() ----- 13: 14: 15: 16: 17: 18: 19: ----- > update phase 20: prev.next = new 21: new.lock() 22: new.next = succ 23: 24: prev.unlock() 25: 26: new.unlock() 27: 28: succ.unlock() </pre>	<pre> 1: FUNCTION addThird(List list, Node new) ----- > read-only phase 2: lockedSet.init(), readSet.init() 3: if !track(list) then goto 1 4: Node prev = list.head 5: if !track(prev) then goto 1 6: lockedSet.remove(list) 7: Node succ = prev.next 8: if !track(succ) then goto 1 9: lockedSet.remove(prev) 10: prev = succ 11: succ = succ.next 12: if !track(succ) then goto 1 ----- > validation phase 13: for all obj in lockedSet do 14: if !obj.tryLock() then 15: unlockAll() 16: goto 1 17: if !validateReadSet() then 18: unlockAll() 19: goto 1 ----- > update phase 20: prev.next = new 21: new.lock() 22: new.next = succ 23: prev.version++ 24: prev.unlock() 25: new.version++ 26: new.unlock() 27: succ.version++ 28: succ.unlock() </pre>
(a) Code with original locking	(b) The code produced by our transformation

Figure 1: Code example. The synchronization code is in bold.

```

1: FUNCTION track(obj)
2:   lockedSet.add(obj)
3:   long ver = obj.version
4:   readSet.add((obj,ver))
5:   if (obj,) ∈ readSet and v!=ver then
6:     return false
7:   if obj.isLockedByAnotherThread() then
8:     return false
9:   return true

```

Figure 2: In read-only phase, locking is replaced by tracking locks and read objects’ versions.

```

1: FUNCTION validateReadSet()
2:   for all (obj,ver) in readSet do
3:     if obj.isLockedByAnotherThread()
4:       then
5:         return false > validation failed (locked
6:         object)
7:         if obj.version != ver then
8:           return false > validation failed
9:           (different version)
10:    retrun true > validation succeed

```

Figure 3: Read set validation.

struction. If a tryLock fails, the code unlocks all previously acquired locks and restarts from the beginning (lines 13-16).

The function *validateReadSet* in Figure 3 is used to validate past reads: it returns *true* if and only if the objects in the read set have not been updated. The function checks that each object in the read set is not locked by another thread, and that the object’s current version matches the version saved in the *readSet*. This check guarantees that the object was not locked from the time it was read until the time it was validated. Since operations write only to locked objects, it follows that the object was not changed. This *readSet* validation can be viewed as a double collect [4] of all objects accessed by the read-only phase. The operation is restarted if the validation fails (lines 17-19).

Update Phase.

In this phase our transformation preserves the original locking while maintaining the versions of the objects, i.e.,

the version of an object *o* is incremented every time *o* is unlocked. Here, before each unlock instruction *x.unlock()* we insert the code *x.version++*. An example is shown in lines 20-28 of Figure 1b.

2.3.2 Exceptions from Regular Flow

Other than in Figure 2, the read phase does not validate past reads during its executions — as a result, it may observe an inconsistent state of shared memory (as explained, e.g., in [27]).

In order to avoid infinite loops (in the read-only phase) that might occur due to inconsistent reads, a timeout is set. If the timeout expires before the read-only phase is completed, read set validation takes place (by invoking the function *validateReadSet*). If the validation fails, the operation is restarted. This is realized by inserting code that examines the timeout in every loop iteration in the original code.

Similarly, inconsistent views may lead to spurious exceptions in the read-only phase. We therefore catch all exceptions, and perform validation. Here too, if the validation fails, the operation is restarted. Otherwise, the exception is handled as in the original code.

Note that, using our transformation, the shared state at the end of the validation phase is identical to the state that would have been reached had the code been executed pessimistically from the outset. Hence, the three-phase version of the code is compatible with the instrumented pessimistic version. This means that if the optimistic phase is unsuccessful for any reason, we can always fall back on the pessimistic version. Moreover, we can switch from optimistic to pessimistic synchronization *at any point* during the read phase. We use this property in two ways, as we now describe.

First, we avoid livelocks by limiting the number of restarts due to conflicts: The validation phase tracks the number of restarts in a thread-local variable. If this number exceeds a certain threshold, we perform the entire operation optimistically.

Second, this property offers the optimistic implementation the liberty of failing spuriously, even in the absence of conflicts, because it can always fall back on the safe pessimistic version of the code. We take advantage of this liberty, and implement the *lockedSet* and *readSet* by using constant size arrays. In case either of these arrays becomes full, we cannot proceed with the optimistic version, but also do not need to start the operation anew. Instead, we immediately perform the validation phase, which, if successful, switches to a pessimistic modus operandi, after having acquired all the needed locks.

3. ANALYSIS

We argue that our transformation is *correct*, in the sense that all the external behaviors of the synthesized code are allowed by the original implementation. This implies that the transformed code preserves essential properties of the original, such as serializability, linearizability, and deadlock-freedom. In this section we provide informal correctness arguments; a formal proof is deferred to Appendix A.

Indistinguishability of successful executions.

Every operation execution of the modified code starts at the beginning of the read-only phase and, (barring exceptions, deadlocks, or infinite loops, which will be discussed later), ends in a return statement in the update phase. To show the correctness of such an execution, we argue that it is indistinguishable from an execution of the entire operation using the update phase. Since the instrumented code maintains version numbers that are not present in the original code, we need to project these out of the state in order to argue that the same states are reached in both cases. For a state s (i.e., an assignment of values to shared and local variables) in an execution of the instrumented code, we denote by \hat{s} the same state without version numbers.

The indistinguishability hinges on the following properties:

1. *The read-only phase has no side effects* since it does not modify shared state.
2. *An unsuccessful validation phase has no lasting effect on shared state*, since it releases all locks it acquires. Moreover, it does not cause deadlock thanks to the

use of tryLocks and releasing all locks upon failure to acquire one. Since we assume that the update phase only uses blocking lock calls, the only impact an unsuccessful validation phase can have on it is additional waiting on locks, which does not impact its external behavior.

3. *A successful validation phase executed after point ℓ in the read-only phase leads to a memory state s , such that \hat{s} would have been reached by performing the original update phase (from the beginning) until point ℓ .*

While the first two properties can be directly observed from the code, the third requires more careful reasoning. To understand why it is correct, recall that, in the read-only phase, the only difference between the optimistic execution and the original one is that reads are performed without holding a lock. We argue that the execution is equivalent to one that would have acquired and released locks as in the code of the update phase. To this end, we show that the object was unlocked (and hence, was not modified), for the duration of entire the period when the original code would have locked it. We distinguish between two cases:

1. In case the lock would have been released by original code before location ℓ , we argue that the lock was free during the read-only phase from the time when it would have been locked and at least until the time when it would have been released by the original code. This is because the successful validation means that (1) the object is currently unlocked, and (2) object's version had not increased from the time it was added to the *readSet* in the read-only phase (i.e., the *lock* instruction in the original code), and until the validation phase, i.e., after it would have been released. Recall that in our instrumented code, the object's version increases every time it is locked.
2. In case the lock would have been held by the original code at location ℓ , it is included in the *lockedSet* during the validation phase, and hence acquired before branching to the update phase. As in the previous case, the successful validation of the *readSet* ensures that the lock had been available since the time when it would have been locked by the original code.

Progress.

Finally, we argue that our read-only and validation phases do not generate spurious deadlocks, exceptions, or infinite loops that were not present in the original code. In other words, if the original code would have successfully completed with a return step, so does the instrumented code.

First, consider deadlocks. The read and validation phases of our instrumented code do not use blocking locks – the read-phase does not use locks at all, whereas the validation phase uses tryLocks. Therefore, both phases are non-blocking. In principle, the optimistic approach may lead to livelocks, but our algorithm falls back on the pessimistic approach following a bounded number of restarts, and hence cannot livelock. We get that any lack of progress must be due to blocking in the update phase. Since this section of the code is unchanged, and since we ensure that it begins when holding the same locks as in the original protocol, we

get that our transformation does not introduce any source of spurious blocking that is not present in the original code.

Exceptions and infinite loops may be introduced in the read-only phase due to reading an inconsistent view of shared memory. Nevertheless, we detect these situations by performing validation both periodically and on exceptions: As we have argued above, the validation phase is successful only if the execution is equivalent to a lock-based one. Thus, every inconsistent view that might lead to a spurious exception or infinite loop causes the validation to fail, and we detect these cases before they have any external impact.

4. EVALUATION

We evaluate the performance of our approach on two types of data structures. In Section 4.1 we consider search trees supporting insert, delete, and get operations, whereas Section 4.2 focuses on data structures that, in addition, support range queries that retrieve all records within a given range. We compare the performance of our approach in terms of throughput to fully pessimistic solutions applying fine-grain locking. These algorithms also serve as the lock-based reference implementation at the base of our semi-optimistic implementations. We further compare our approach to hand-crafted state-of-the-art data structure implementations supporting the same functionality.

We also measured the performance of global lock-based implementations. In all workloads, the results were identical or inferior to those achieved by pessimistic fine-grain locking. We hence omitted these results to avoid obscuring the presentation.

We follow a standard evaluation methodology ([18, 40, 12, 8]). Each experiment consists of 5 trials. A trial is a five second run in which each thread continuously executes randomly chosen operations drawn from the workload distribution, with keys selected uniformly at random from the range $[0, 2 \cdot 10^6]$. Each trial begins by initiating a new data structure and applying an untimed pre-filling phase, which continues until the size of the data structure is within 5% of 10^6 records. The presented results are the average throughput over all trials (with the pre-filling phase excluded). We also experimented with a smaller range ($[0, 2 \cdot 10^4]$) to test different contention levels; since the results showed similar trends, they are omitted here.

All implementations are written in Java. We ran the experiments on a dedicated machine with four Intel Xeon E5-4650 processors, each with 8 cores, for a total of 32 threads (with hyper-threading disabled). We used Ubuntu 12.04.4 LTS and Java Runtime Environment (build 1.7.0_51-b13) using the 64-Bit Server VM (build 24.51-b03, mixed mode).

4.1 Insert-Delete-Get Operations

We start by benchmarking a search-tree supporting the basic insert, delete, and get (lookup) operations. Our experiments evaluate unbalanced as well as balanced trees.

We employ textbook sequential implementations of an unbalanced binary tree, and a treap [7]. To generate pessimistic lock-based implementations, we synthesize concurrent code by applying the domination locking technique to the sequential data structures. The resulting algorithms are denoted Lock-Tree and Lock-Treap. Finally, we apply our lock-removal transformation to the reference implementations to get our semi-optimistic versions of the code, which we call Auto-Tree and Auto-Treap, respectively.

We further compare our implementations to their hand-crafted state-of-the-art counterparts. We compare Auto-Tree to

LO-Tree The locked-based unbalanced tree of Drachslers et al. [18]².

Auto-Treap is evaluated against three hand-crafted implementations

LO-AVL The locked-based relaxed balanced AVL tree of Drachslers et al. [18]³.

Snap-Tree The locked based relaxed balanced AVL tree of Bronson et al. [11]⁴.

Java-Skiplist The non-blocking skip-list in the the Java standard library; based on the work of Fraser and Harris [23].

We evaluate performance in three representative workloads distributions: a *read-only* workload comprised of 100% lookup operations, a *write-dominated* workload consisting of insert and delete operations (50% each), and a *mixed workload* with 50% lookups, 25% inserts, and 25% deletes.

Results.

Figure 4 shows the throughput of unbalanced data structures and Figure 5 shows the throughput of the balanced ones. We see that our semi-optimistic solution is far superior to previous, fully-pessimistic, automated approaches, with 50x to 70x higher performance. It successfully overcomes the bottlenecks associated with lock contention in the Lock-Tree and Lock-Treap implementations. In many scenarios, our solution comes close to custom-tailored implementations.

The results for the read-only workload show the main overhead of our automatic approach. By profiling the code, we learned that the bulk of this overhead stems from the need to track all read objects, which is inherent to our transformation. This is in contrast with the hand-crafted implementations, which have small overhead on reads in this scenario, thanks to either wait-free reads (in LO-Tree and LO-AVL), or optimistic validation (in Snap-Tree).

As the ratio of updates in the workload increases, our automatic implementation closes this gap. In other words, the transformed code deals well with update contention. This might be due to the fact that once an update phase begins, the operation is not delayed due to concurrent read-only operations.

4.2 Range Queries

Next we evaluate the performance of our approach when the data structure supports a more intricate functionality like range queries. We use a skip list, which readily supports range queries by nature of its linked-structure. The core of the implementation is the key lookup method; once reaching the key, a record can be added or be removed in place, and an iteration of subsequent records can be executed by traversing the bottom-level linked-list.

The domination locking scheme cannot be efficiently applied to the skip list structure since it is a DAG rather than

²Implementation provided by the authors.

³Implementation available at <https://github.com/logicalordering/trees>

⁴Implementation available at <https://github.com/nbronson/snaptree>




Figure 4: Throughput of unbalanced data structures.




Figure 5: Throughput of balanced data structures.

a tree. Instead, we manually devise a pessimistic locking protocol. Our algorithm, (inspired by the one in [30]), applies hand-over-hand locking at each level, so that at the end of the search, the operation holds locks on two records in each level, which define the minimal interval within this level containing the lookup key (or the first lookup key in the case of a range query). Upon reaching the bottom level, unnecessary locks are released, as follows: update operations only keep locks on nodes they intend to modify, whereas range queries keep the locks in the level with the minimal interval spanning the range. Range queries then continue to use hand-over-hand locking to traverse through all records within the range. The use of hand-over-hand locking ensures that range queries are atomic (linearizable), i.e., return a consistent view of the data structure.

This pessimistic lock-based algorithm is denoted Lock-Skiplist. As in previous data structures, we apply the lock-removal transformation to the reference implementation to get a semi-optimistic algorithm, which we call Auto-Skiplist.

Auto-Skiplist is compared to Lock-Skiplist as well as those of the aforementioned state-of-the-art data structures that support range queries. Specifically, we compare our results to (1) Snap-Tree, which provides atomic range queries by traversing a clone of the original tree that is lazily generated, and (2) iterations in Java’s concurrent skip list, which are *not* guaranteed to be atomic. Like many data structure libraries, LO-Tree and LO-AVL do not support atomic range queries, and there is no straightforward way to add them.

The evaluation of range queries focuses on a mixed workload, where half the threads are dedicated to performing range queries, and the other half performs a mix of insert and delete operations (50% each). This mix allows us to evaluate both the performance of the range queries themselves, and their impact on concurrent updates. We experiment with queries with big ranges varying between 1000 to 2000 keys (Figure 6) and small ranges between 10 to 20 keys (Figure 7).

Results.

We measure the throughput of range queries and update operations separately. The overall number of range queries executed per second is reported in Figures 6a and 7a, and the overall number of update (insert and delete) operations executed per second is reported in Figures 6b and 7b.

Again, we see that our transformed code is superior to the fully-pessimistic fine-grain one, improving throughput by up to two orders of magnitude – 75x for small range queries, 4x for big ranges, and 100x for update operations, regardless of the range size. We believe that, as in the domination locking versions of the tree data structures, holding a lock on the head sentinel of the skip list in Lock-Skiplist, even for short periods, imposes a major performance penalty, which is eliminated by our semi-optimistic approach.

The performance of Auto-Skiplist is almost identical to that of Java-Skiplist, despite the fact that Java-Skiplist performs inconsistent iterations.

Figure 6a shows that, for big ranges, Snap-Tree scales well up to 8 threads, outperforming all other implementations, but at 32 threads its performance deteriorates. This might be because this implementation is optimized for full scans or very big range queries running sequentially. The overhead of initiating a clone of the data structure per range query, which involves waiting for all pending update operations to

complete, hampers scalability when a number of queries are executed in parallel. This effect is even more pronounced when the ranges are small (Figure 7a), and the throughput of Snap-Tree flattens out.

In addition, the lazy cloning required to support range queries imposes an overhead on update operations. Copying each node during downward traversal might be the main impediment preventing these operations from scaling when the ranges are small (Figure 7b) as well as when they are big (Figure 6b). In contrast, the update operations in Auto-Skiplist and Java-Skiplist do not take any special measures for the benefit of concurrent queries (which validate themselves in Auto-Skiplist, and are not atomic in Java-Skiplist), and hence continue to perform well in their presence.

5. RELATED WORK

Concurrent Data Structures.

Many sophisticated concurrent data structures (e.g., [8, 18, 40, 12, 15, 9, 5, 20, 11, 29, 23, 38]) were developed and used in concurrent software systems [44]. Implementing efficient synchronization for such data structures is considered a challenging and error-prone task [44, 17, 33]. As a result, concurrent data structures are manually implemented by concurrency experts. This paper shows that (in some cases) an automatic algorithm can produce synchronization that is comparable to synchronization implemented by experts.

Lock Inference Algorithms.

There has been a lot of work on automatically inferring locks for transactions. Most algorithms in the literature infer locks that follow the two-phase locking protocol [36, 21, 25, 14, 31, 16]. Our approach can potentially be used to optimize the synchronization produced by these algorithms. For example, for algorithms that employ a two-phase variant in which all locks are acquired at the beginning of a transaction (e.g., [25, 14]), our approach may be used to defer the locking (e.g., to just before the first write operation) and even to eliminate some of the locking steps. We demonstrate the benefit of combining our transformation with such algorithms by using the domination locking protocol [24] to produce efficient concurrency control for dynamic data structures.

Transactional Memory.

Transactional memory approaches (TMs) dynamically resolve inconsistencies and deadlocks by rolling back partially completed transactions. Unfortunately, in spite of a lot of effort and many TM implementations (see [27]), existing TMs have not been widely adopted due to various concerns [19, 13, 37], including high runtime overhead, poor performance and limited ability to handle irreversible operations. Modern concurrent programs and data structures are typically based on hand-crafted synchronization, rather than on a TM approach [44].

Lock Elision.

Our transformation is inspired by the idea of *sequential locks* [37] and the approach presented in [39], which replace locks with optimistic concurrency control in read-only transactions. But in contrast to these works, we handle read-

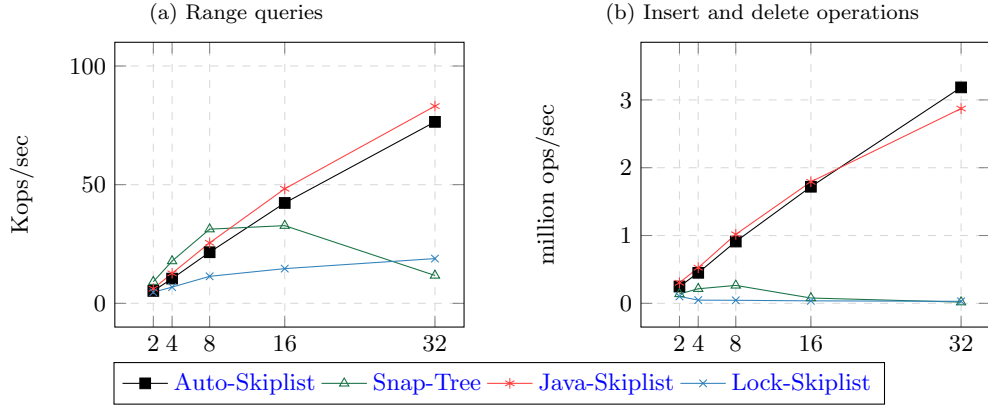


Figure 6: Half the threads execute big range queries $[1000, 2000]$ and half the threads execute insert and delete operations.

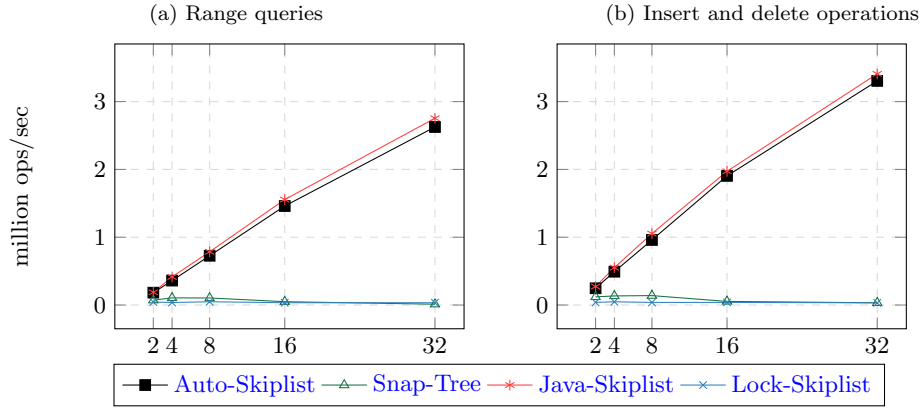


Figure 7: Half the threads execute small range queries $[10, 20]$ and half the threads execute insert and delete operations.

only prefixes of transactions (operations) that do update the shared memory. In fact, as shown in Section 4, our approach is best suited for update-dominated workloads. Moreover, using these approaches for a highly-contended data structure (as in Section 4) is likely to provide limited performance, because each update transaction causes many read-only transactions to abort.

Other works have proposed using transactional memory in order to elide locks from arbitrary critical sections, and fall back on lock-based execution in cases of aborts (e.g., [41, 42, 6]). In contrast to our approach, however, lock elision does not combine speculative and non-speculative execution within the same transaction.

6. DISCUSSION

The development of scalable concurrent programs today heavily relies on custom-tailored implementations, which require painstaking correctness proofs. In this paper, we have shown a relatively simple transformation that can facilitate this labor-intensive process, and thus make scalable synchronization more readily available. The input for our transformation is a conventional lock-based concurrent program, which may be either constructed manually or synthesized from sequential code. Our source-to-source transformation then makes judicious use of optimism in order to eliminate principal concurrency bottlenecks in the given program and improve its scalability.

We have illustrated our method for a number of data structures – unbalanced and balanced search trees, as well as skip lists supporting range queries. In all cases, the transformed code performed significantly better than the original lock-based one. It also scaled comparably to hand-crafted implementations that took considerably more effort to produce. Moreover, extending the synthesized code with new functionalities such as range queries was immediate. In these examples, we have manually applied our transformation (according to the algorithm presented in Section 2). An interesting direction for future work would be to create a tool that automatically applies our transformation at compile time.

Our approach makes use of a common pattern in data structures, where an operation typically begins with a long read-only traversal, followed by a handful of (usually local) modifications. A promising direction for future work is to try and exploit similar patterns in order to parallelize or remove locks in other types of code (not data structures), for example, programs that rely on two-phase locking. Furthermore, for programs that follow different patterns, other combinations of optimism and pessimism may prove effective.

Finally, there still remains a gap between the performance achievable by manually optimized solutions and what we could achieve automatically. Our algorithm induces inherent overhead for tracking all operations in the read-only phase for later verification. In specific data structures, these checks might be redundant, but it is difficult to detect this automatically. We believe that it may well be possible to enhance transformations such as ours with computer-assisted optimizations. For example, a programmer may provide hints regarding certain invariants that are always preserved in the code, in order to eliminate the need for tracking some values for later validation. Such optimizations have the potential to bridge the remaining performance gap, while requiring far less work for proving correctness – instead of proving

that the entire construction is correct, the developer would only need to prove that her program maintains the specific invariants used.

7. REFERENCES

- [1] Concurrentskiplistmap from java.util.concurrent. docs.oracle.com/javase/7/docs/api/java/util/concurrent/ConcurrentSkipListMap.html.
- [2] A fast and lightweight key/value database library by google. <http://code.google.com/p/leveldb>.
- [3] jmonkeyengine: a 3d game engine for java developers. <http://jmonkeyengine.org/>.
- [4] Y. Afek, H. Attiya, D. Dolev, E. Gafni, M. Merritt, and N. Shavit. Atomic snapshots of shared memory. *J. ACM*, 40(4):873–890, Sept. 1993.
- [5] Y. Afek, H. Kaplan, B. Korenfeld, A. Morrison, and R. E. Tarjan. CBTree: A practical concurrent self-adjusting search tree. In *DISC*, pages 1–15, 2012.
- [6] Y. Afek, A. Levy, and A. Morrison. Software-improved hardware lock elision. In *PODC*, 2014.
- [7] C. R. Aragon and R. Seidel. Randomized search trees. In *30th Annual Symposium on Foundations of Computer Science, Research Triangle Park, North Carolina, USA, 30 October - 1 November 1989*, pages 540–545, 1989.
- [8] M. Arbel and H. Attiya. Concurrent updates with RCU: search tree as an example. In *PODC*, pages 196–205, 2014.
- [9] A. Braginsky and E. Petrank. A lock-free B+tree. In *SPAA*, pages 58–67, 2012.
- [10] N. Bronson. *Composable Operations on High-Performance Concurrent Collections*. PhD thesis, Stanford University, Dec. 2011.
- [11] N. G. Bronson, J. Casper, H. Chafi, and K. Olukotun. A practical concurrent binary search tree. In *PPOPP*, pages 257–268, 2010.
- [12] T. Brown, F. Ellen, and E. Ruppert. A general technique for non-blocking trees. In *19th ACM Symposium on Principles and Practice of Parallel Programming (PPoPP)*, pages 329–342, 2014.
- [13] C. Cascaval, C. Blundell, M. Michael, H. W. Cain, P. Wu, S. Chiras, and S. Chatterjee. Software transactional memory: Why is it only a research toy? *Queue*, 6(5):46–58, Sept. 2008.
- [14] S. Cherem, T. Chilimbi, and S. Gulwani. Inferring locks for atomic sections. In *PLDI*, 2008.
- [15] T. Crain, V. Gramoli, and M. Raynal. A contention-friendly binary search tree. In *19th International Conference on Parallel Processing (Euro-Par)*, pages 229–240, 2013.
- [16] D. Cunningham, K. Gudka, and S. Eisenbach. Keep off the grass: Locking the right path for atomicity. In *CC*, pages 276–290, 2008.
- [17] S. Doherty, D. L. Detlefs, L. Groves, C. H. Flood, V. Luchangco, P. A. Martin, M. Moir, N. Shavit, and G. L. Steele, Jr. Dcas is not a silver bullet for nonblocking algorithm design. In *SPAA*, 2004.
- [18] D. Drachslar, M. T. Vechev, and E. Yahav. Practical concurrent binary search trees via logical ordering. In *PPoPP*, pages 343–356, 2014.

- [19] J. Duffy. A (brief) retrospective on transactional memory. 2010.
<http://joeduffyblog.com/2010/01/03/\a-brief-retrospective-on-transactional-memory/>.
- [20] F. Ellen, P. Fatourou, E. Ruppert, and F. van Breugel. Non-blocking binary search trees. In *29th ACM Symposium on Principles of Distributed Computing (PODC)*, pages 131–140, 2010.
- [21] M. Emmi, J. S. Fischer, R. Jhala, and R. Majumdar. Lock allocation. In *POPL*, pages 291–296, 2007.
- [22] D. G. Ferro, F. Junqueira, I. Kelly, B. Reed, and M. Yabandeh. Omid: Lock-free transactional support for distributed data stores. In *IEEE 30th International Conference on Data Engineering, Chicago, ICDE 2014, IL, USA, March 31 - April 4, 2014*, pages 676–687, 2014.
- [23] K. Fraser. *Practical lock-freedom*. PhD thesis, University of Cambridge, 2004.
- [24] G. Golan-Gueta, N. G. Bronson, A. Aiken, G. Ramalingam, M. Sagiv, and E. Yahav. Automatic fine-grain locking using shape properties. In *OOPSLA*, pages 225–242, 2011.
- [25] K. Gudka, T. Harris, and S. Eisenbach. Lock inference in the presence of large libraries. In *ECOOP*. 2012.
- [26] R. Guerraoui and M. Kapalka. On the correctness of transactional memory. In *PPOPP*, pages 175–184, 2008.
- [27] T. Harris, J. Larus, and R. Rajwar. Transactional memory, 2nd edition. *Synthesis Lectures on Computer Architecture*, 5(1), 2010.
- [28] M. Herlihy and E. Koskinen. Transactional boosting: a methodology for highly-concurrent transactional objects. In *PPoPP*, pages 207–216, 2008.
- [29] M. Herlihy, Y. Lev, V. Luchangco, and N. Shavit. A simple optimistic skiplist algorithm. In *14th International Conference on Structural Information and Communication Complexity (SIROCCO)*, pages 124–138, 2007.
- [30] M. Herlihy and N. Shavit. *The Art of Multiprocessor Programming*. Morgan Kaufmann Publishers Inc., 2008.
- [31] M. Hicks, J. S. Foster, and P. Prattikakis. Lock inference for atomic sections. In *Proceedings of the First ACM SIGPLAN Workshop on Languages, Compilers, and Hardware Support for Transactional Computing*, June 2006.
- [32] A. Israeli and L. Rappoport. Disjoint-access-parallel implementations of strong shared memory primitives. In *Proceedings of the Thirteenth Annual ACM Symposium on Principles of Distributed Computing, PODC '94*, pages 151–160, New York, NY, USA, 1994. ACM.
- [33] G. Jin, W. Zhang, D. Deng, B. Liblit, and S. Lu. Automated concurrency-bug fixing. In *OSDI*, 2012.
- [34] E. Koskinen, M. Parkinson, and M. Herlihy. Coarse-grained transactions. In *POPL*, pages 19–30, 2010.
- [35] K. Lev-Ari, G. Chockler, and I. Keidar. On correctness of data structures under reads-write concurrency. In *DISC*, October 2014.
- [36] B. McCloskey, F. Zhou, D. Gay, and E. Brewer. Autolocker: synchronization inference for atomic sections. In *POPL*, pages 346–358, 2006.
- [37] P. E. McKenney. Is parallel programming hard, and, if so, what can you do about it? *Linux Technology Center, IBM Beaverton*, Aug. 2012.
- [38] M. M. Michael and M. L. Scott. Simple, fast, and practical non-blocking and blocking concurrent queue algorithms. In *PODC*, 1996.
- [39] T. Nakaike and M. M. Michael. Lock elision for read-only critical sections in java. In *PLDI*, pages 269–278, 2010.
- [40] A. Natarajan and N. Mittal. Fast concurrent lock-free binary search trees. In *19th ACM Symposium on Principles and Practice of Parallel Programming (PPoPP)*, pages 317–328, 2014.
- [41] R. Rajwar and J. R. Goodman. Transactional lock-free execution of lock-based programs. *SIGOPS Oper. Syst. Rev.*, 36(5):5–17, Oct. 2002.
- [42] A. Roy, S. Hand, and T. Harris. A runtime system for software lock elision. In *EuroSys*, 2009.
- [43] O. Shacham. *Verifying Atomicity of Composed Concurrent Operations*. PhD thesis, Tel Aviv University, 2012.
- [44] O. Shacham, N. Bronson, A. Aiken, M. Sagiv, M. Vechev, and E. Yahav. Testing atomicity of composed concurrent operations. In *OOPSLA*, 2011.
- [45] O. Shalev and N. Shavit. Predictive log-synchronization. In Y. Berbers and W. Zwaenepoel, editors, *Proceedings of the 2006 EuroSys Conference, Leuven, Belgium, April 18-21, 2006*, pages 305–315. ACM, 2006.
- [46] F. Zulkaryarov, V. Gajinov, O. S. Unsal, A. Cristal, E. Ayguadé, T. Harris, and M. Valero. Atomic quake: using transactional memory in an interactive multiplayer game server. In *ACM Sigplan Notices*, volume 44, pages 25–34. ACM, 2009.

APPENDIX

A. FORMAL CORRECTNESS PROOF

We now formalize the correctness arguments made in Section 3. First we define our model and the correctness properties of the algorithm for which we provide the proof.

Model.

We consider an asynchronous shared memory model, where independent threads interact via shared memory objects. Every thread executes a sequence of operations, each of which is invoked with certain parameters and returns a response. An operation’s execution consists of a sequence of primitive *steps*, beginning with an *invoke* step, followed by atomic accesses to shared objects, and ending with a *return* step. Steps also modify the executing thread’s local variables.

A *configuration* is an assignment of values to all shared and local variables. Thus, each step takes the system from one configuration to another. Steps are deterministically defined by the data structure’s protocol and the current configuration. In the *initial configuration*, each variable holds its initial value.

An *execution* is an alternating sequence of configurations and steps, $C_0, s_1, C_1, \dots, s_i, C_i, \dots$, where C_0 is an initial

configuration, and each configuration C_i is the result of executing step s_i on configuration C_{i-1} . We only consider finite executions in this paper. An execution is *sequential* if steps of different operations are not interleaved. In other words, a sequential execution is a sequence of operation executions.

Two executions are *indistinguishable* to a set of operations if each operation in the set executes the same steps on primitive shared objects, and receives the same value from those primitives, in both executions. A step τ by operation op is *invisible* to all other operations if the executions with and without τ are indistinguishable to $OP \setminus \{op\}$. For example, read steps are invisible.

Correctness.

The correctness of a data structure is defined in terms of its external behavior, as reflected in values returned by invoked operations. Correctness of a code transformation is proven by showing that the synthesized code's executions are equivalent to ones of the original code, where two executions are *equivalent* if every thread invokes the same operations in the same order in both executions, and gets the same result for each operation. More formally, we say in this paper that a code transformation is *correct* if every execution of the transformed code is equivalent to some execution of the original code.

The widely-used correctness criterion of serializability relies on equivalence to sequential executions in order to link a data structure's behavior under concurrency to its sequentially specified behavior. Since equivalence is transitive, we get that any code transformation satisfying our correctness notion, when applied to serializable code, yields code that is also serializable. If the code transformation further ensures the real-time order of operations (i.e., operations that do not overlap appear in the same order in executions of the transformed and original code), then linearizability (atomicity) is also invariant under the transformation. Another important aspect of correctness is preserving the progress conditions of the original code, for example, deadlock-freedom.

In this paper, we are not concerned with internal consistency (as required e.g., by opacity [26] or the validity notion of [35]), which restricts the configurations an operation might see during its execution. This is because our code transformation uses timeouts and exception handlers to overcome unexpected behavior that may arise when a thread sees an inconsistent view of global variables (similar to [39]).

Formal Proof.

We consider a finite execution π of the transformed algorithm, and find an equivalent execution of the original lock-based algorithm. Let OP be the set of operations in π . For each operation we consider its *successful validation*, i.e., the successful execution of a validation phase when switching from the read-only phase to the update phase. Each operation executes at most one successful validation. Denote by e_1, e_2, \dots, e_k the sequence of the first steps of the read set validation in the execution of successful validation phases, by their order in π , where e_i is a step of the operation op_i executed by process p_i . (Possibly $p_i = p_j$ for $j \neq i$).

For every operation op_i , consider the partition of π to the following intervals $\pi = \alpha_i \beta_i \gamma_i$, such that α_i includes the execution interval of op_i 's read-only phase (denote op_i 's read set rs_i); $\beta_i = \beta_{i1} \beta_{i2}$, is the minimal execution interval

of op_i 's successful validation phase; in β_{i1} , op_i acquires locks on its lock set, denoted ls_i ; e_i is the first step of β_{i2} , namely the read set validation interval.

The next claim follows from the fact that the validation phase of op_i in β_i is successful, and includes locks and versions re-validation:

CLAIM A.1. *No operation in $OP \setminus \{op_i\}$ holds or acquires a lock in $\alpha_i \beta_{i1}$ on an object obj in rs_i after op_i 's last read of obj in α_i .*

Towards proving equivalence to the original code execution we first project object versions out of π 's configurations, and remove all accesses (reads and writes) to object versions. That is, we replace steps that access versions with local steps that modify the operation's local memory only. (Note that we get an execution with exactly the same invocations, responses, local states, and shared object states, but without versions). Second, note that each operation $op \in OP$ executes at most one successful read-only phase, namely a complete read-only phase followed by a successful validation phase. For each such op , we remove the prefix of op that precedes the successful read-only phase. This includes completely removing operations that have no successful read-only phase. Finally, we remove all validation phases executed during the successful read-only phase that are not the unique successful validation phase of the operation. We call the resulting execution π' .

Essentially, π' is a projection of π excluding versions, all prefixes of the operations preceding their (single) successful read-only phase, and all failed validation attempts. These prefixes and validation phases include read steps as well as tryLock and unlock steps. Removing read steps is invisible to other processes. Since we remove all tryLock steps that have failed to acquire locks, all remaining tryLock are successful also in π' . Therefore, in π' all operations in OP return the same values as in π :

CLAIM A.2. *π and π' are equivalent.*

Note that by construction, in π' every step accessing a lock object, either for locking it or for validating it is not locked, finds the object not locked.

Our main lemma constructs the execution of a fully-pessimistic locking code. The core idea is to replace the optimistic read-only phase and validation phase of each operation with a solo execution of the pessimistic lock-based read phase taking place at the point where the tryLockAll completes.

LEMMA A.3. *There is an execution of lock-based algorithm that is equivalent to π' .*

PROOF. We start with the execution $\pi_0 = \pi'$. For every $i \geq 0$, we show how to perturb π_i to obtain an execution π_{i+1} . For each $j \geq i + 1$, let $\beta'_j = \beta'_{j1} \beta'_{j2}$ be the minimal interval containing op_j 's validation phase in π_{i+1} , where β'_{j1} is the minimal interval containing op_j 's tryLock phase. Denote the configuration between β'_{j1} and β'_{j2} C_j . In π_{i+1} the following conditions are satisfied:

1. The operations op_1, \dots, op_i follow the fully-pessimistic locking algorithm, while the rest of the operations $OPT_{i+1} = OP \setminus \{op_1, \dots, op_i\}$ proceed according to our semi-optimistic algorithm.
2. For $j \geq i + 1$, no operation in $OP \setminus \{op_j\}$ holds a lock in C_j on an object obj in rs_j .

3. For $j \geq i + 1$, no operation in $OP \setminus \{op_j\}$ writes to an object obj in rs_j after op_j last read obj before C_j .
4. For $j \geq i + 1$, all try-lock steps by op_j are invisible to $OP \setminus \{op_j\}$.
5. π' and π_{i+1} are equivalent.

For $OPT_{k+1} = \emptyset$, we get an execution where all operations follow the pessimistic locking algorithm, and by Condition 5 π_{k+1} is equivalent to π' and we are done.

The proof is by induction on i . For the base case we consider the execution $\pi' = \pi_0$. Condition 1 holds since none of the operations in this execution follow the full locking algorithm. Conditions 2 and 3 hold by Claim A.1, and since in π' we only remove read and lock steps, and since accesses to objects (other than versions) are similar in π and π' . Condition 4 holds by the observation above. Condition 5 vacuously holds since π' and π_0 are the same execution.

For the induction step, assume $OPT_i \neq \emptyset$ and the execution $\pi_i = \alpha'_i \beta'_{i_1} \beta'_{i_2} \gamma'_i$ satisfies the above conditions. We replace π_i with $\pi_{i+1} = \alpha''_i \beta''_{i_1} \delta_i \beta''_{i_2} \gamma'_i$, such that α''_i , β''_{i_1} , and β''_{i_2} are the projection of α'_i , β'_{i_1} and β'_{i_2} , excluding the steps by op_i , while δ_i is a p_i -only execution interval in which p_i follows the locking algorithm while reading rs_i ; after δ_i , p_i holds the locks on all objects in ls_i , and holds no lock on other objects. In other words, we replace the optimistic read-only phase and validation phase of op_i with an execution of a read phase instrumented with the locking algorithm, taking place at C_j .

By Condition 2 of the induction hypothesis no operation holds locks on objects in the read set of op_i in C_i , therefore, p_i can acquire the locks on these objects while executing δ_i . By Condition 3 of the induction hypothesis no operation writes to an object obj in rs_i after op_i last read obj before C_i , hence op_i reads the same values in its read set in π_i and π_{i+1} . After δ_i , op_i holds the locks on all objects in ls_i , hence it can continue with the execution of the locking algorithm.

In $\alpha''_i \beta''_{i_1}$ we only removed read steps and tryLock steps by op_i that are invisible to all other operations, by Condition 4. Therefore, the executions $\alpha'_i \beta'_{i_1}$, ending with configuration C' , and $\alpha''_i \beta''_{i_1} \delta_i$, ending with configuration C'' , are indistinguishable to all operations in $OP \setminus \{op_i\}$. In addition, in β''_{i_2} we only removed invisible read steps. The values of all shared objects and locks are the same in C' and C'' , hence the executions $\alpha'_i \beta'_{i_1} \beta'_{i_2} \gamma'_i$ and $\alpha''_i \beta''_{i_1} \delta_i \beta''_{i_2} \gamma'_i$ are indistinguishable to all operations in $OP \setminus \{op_i\}$.

This implies that (1) the projection of the execution π_{i+1} on op_i follows the full locking algorithm satisfying Condition 1, and (2) all operations return the same value in π_{i+1} as in π' , which means Condition 5 holds.

It is left to show that π_{i+1} satisfies Conditions 2, 3, 4. This is straightforward from the induction hypothesis and the fact that only op_i changed its locking pattern in the last iteration, and specifically removed all its try-lock steps, and since δ_i precedes C_j for all $j \geq i + 1$ in π_{i+1} .

□

By Lemma A.3 and Claim A.2 we conclude the following theorem:

THEOREM A.4. *Every execution of the automatic transformation is equivalent to an execution of the locking algorithm.*